
Applied Machine Learning with Big Data

“EE 6973”



Topic:
Convolution Neural Networks

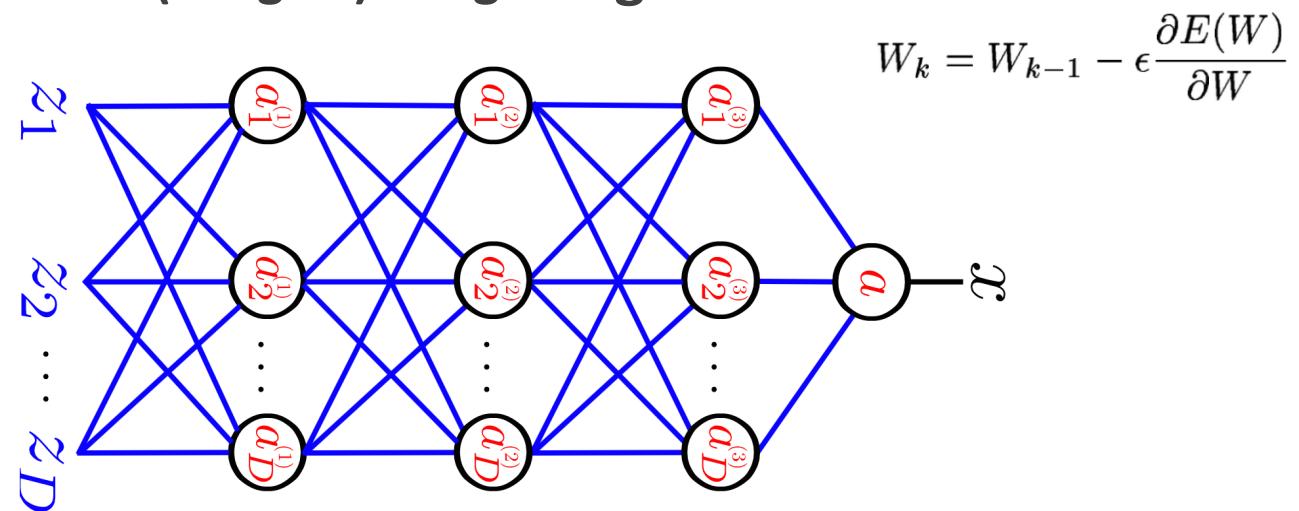
Paul Rad, Ph.D.

Chief Research Officer
UTSA Open Cloud Institute(OCI)
University of Texas at San Antonio

Stochastic Gradient Descent (SGD)

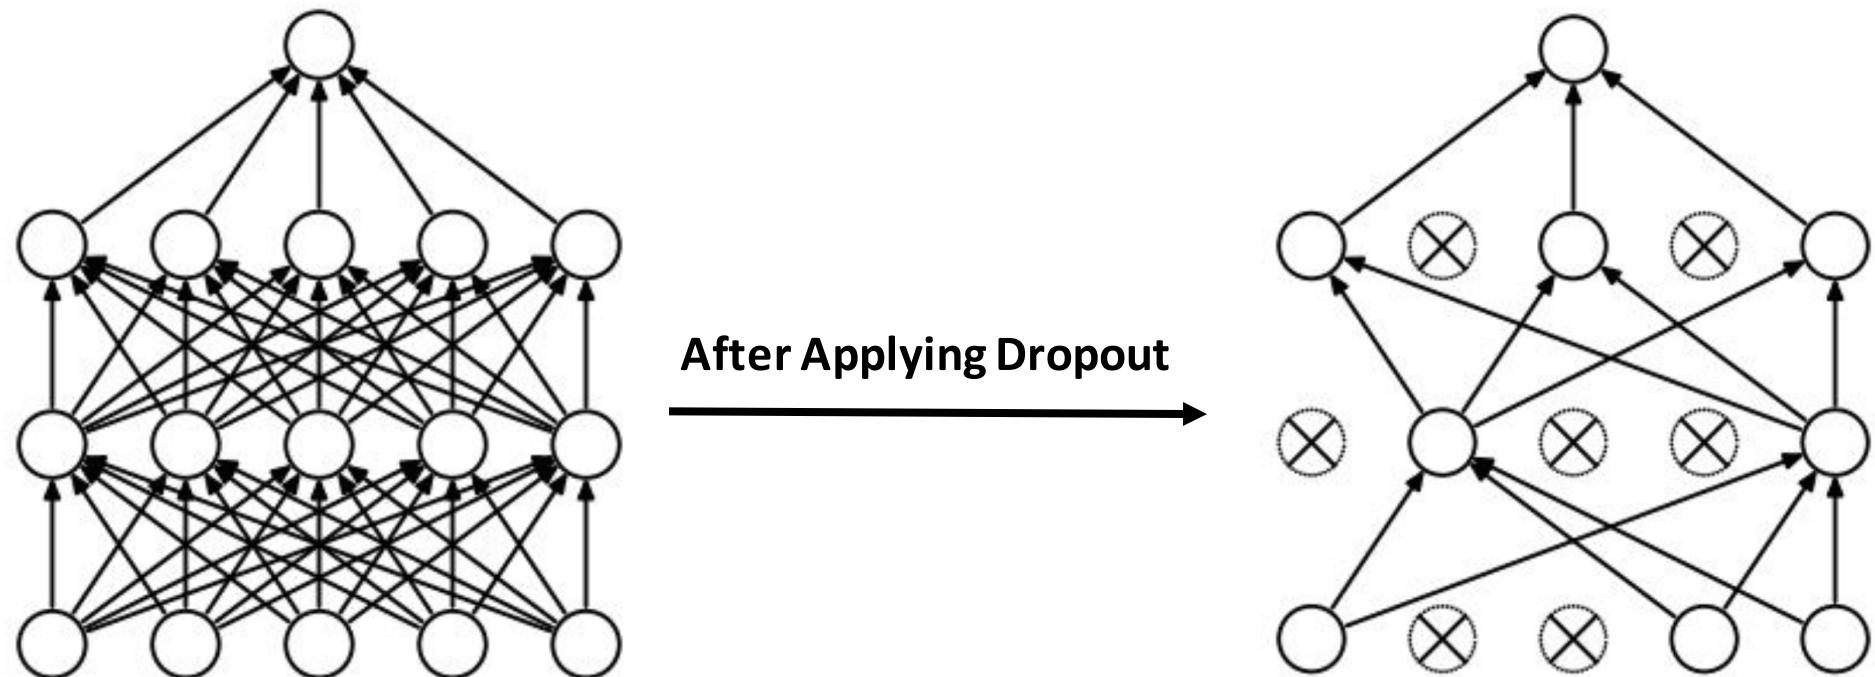
Loop:

1. Sample a batch of data with their labels
2. Forward Propagate it through the graph, calculate the error
3. Backpropagate to calculate the gradients
4. Update the parameters (weights) using the gradient



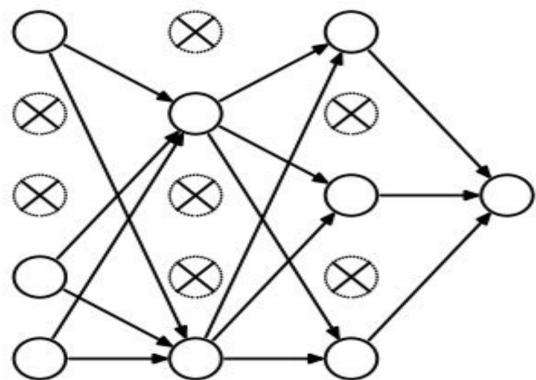
Dropout

Randomly set some neurons to zero in the forward pass



Srivastava et al., 2014, <https://www.cs.toronto.edu/~hinton/absps/JMLRdropout.pdf>

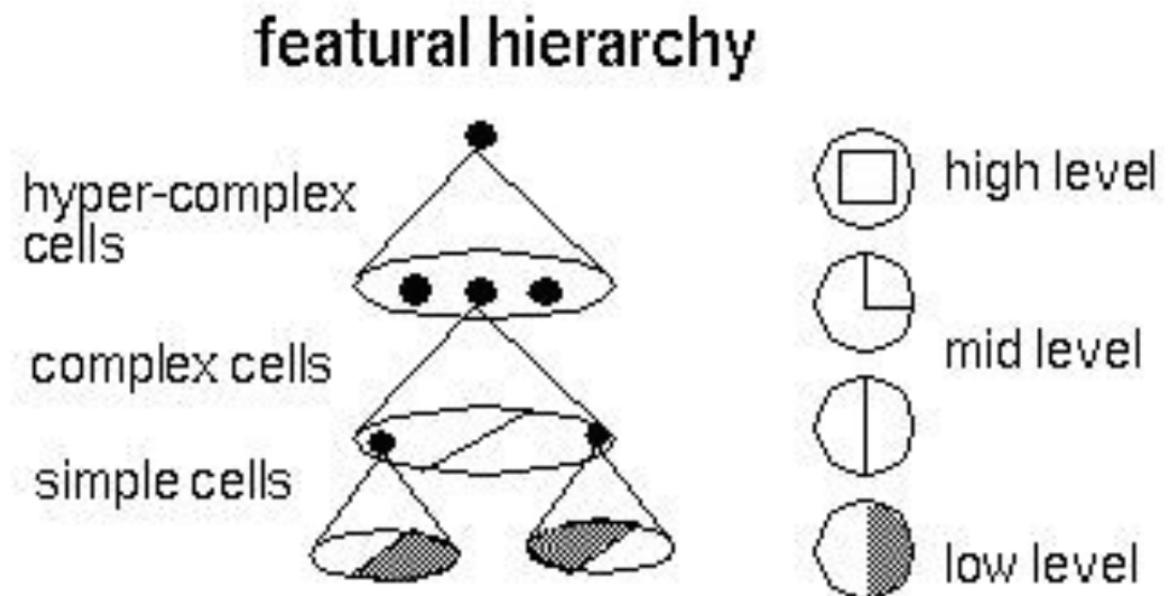
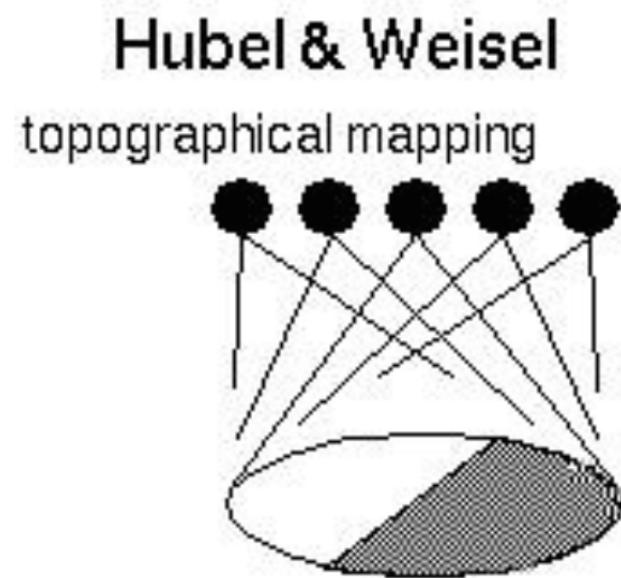
Dropout



**Forces the network to have
a redundant representation**

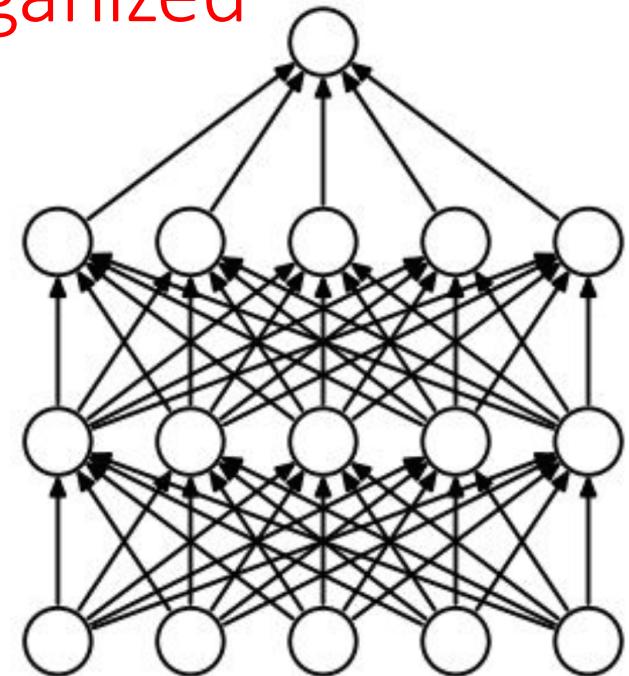
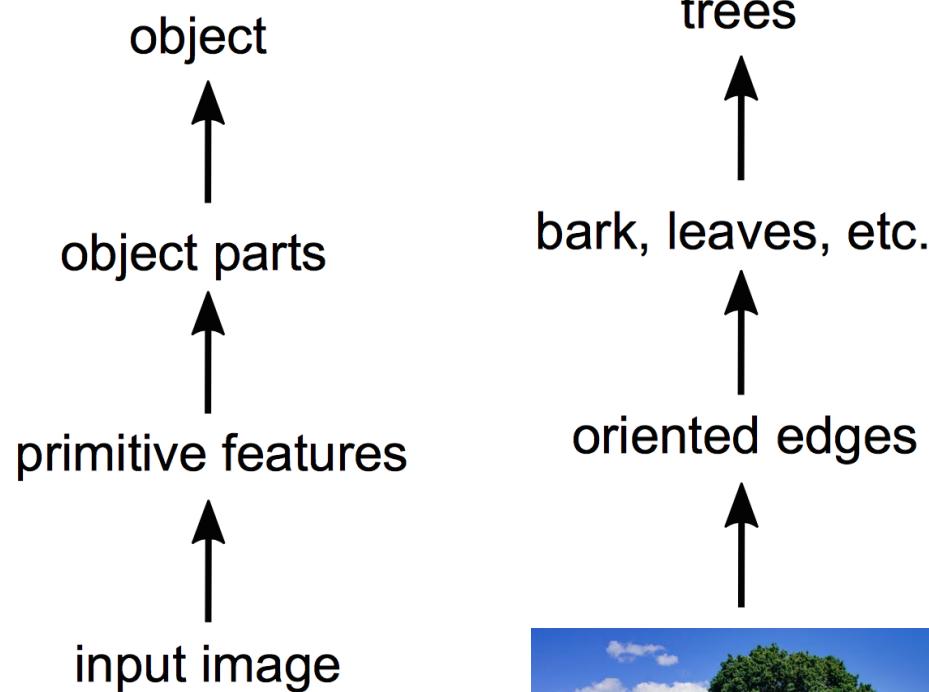


Vision: Hierarchical Organization (Year: 1962)



Why use hierarchical multi-layered models?

Biological vision is hierarchically organized



What's wrong with standard neural networks?

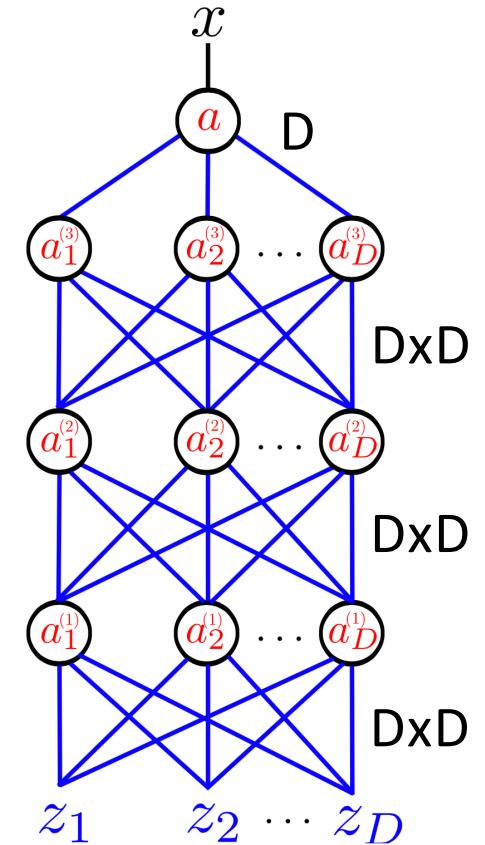
Hard to Train

How many parameters does this network have?

Number of Parameters = $3 \times (D \times D) + D$

For a small $D = 32 \times 32 = 1024$ MNIST image:

Number of Parameters = $3 \times (1024 \times 1024) + 1024$
 $\sim 3 \times 10^6$



Architecture of LeNet-5, Convolution Neural Network

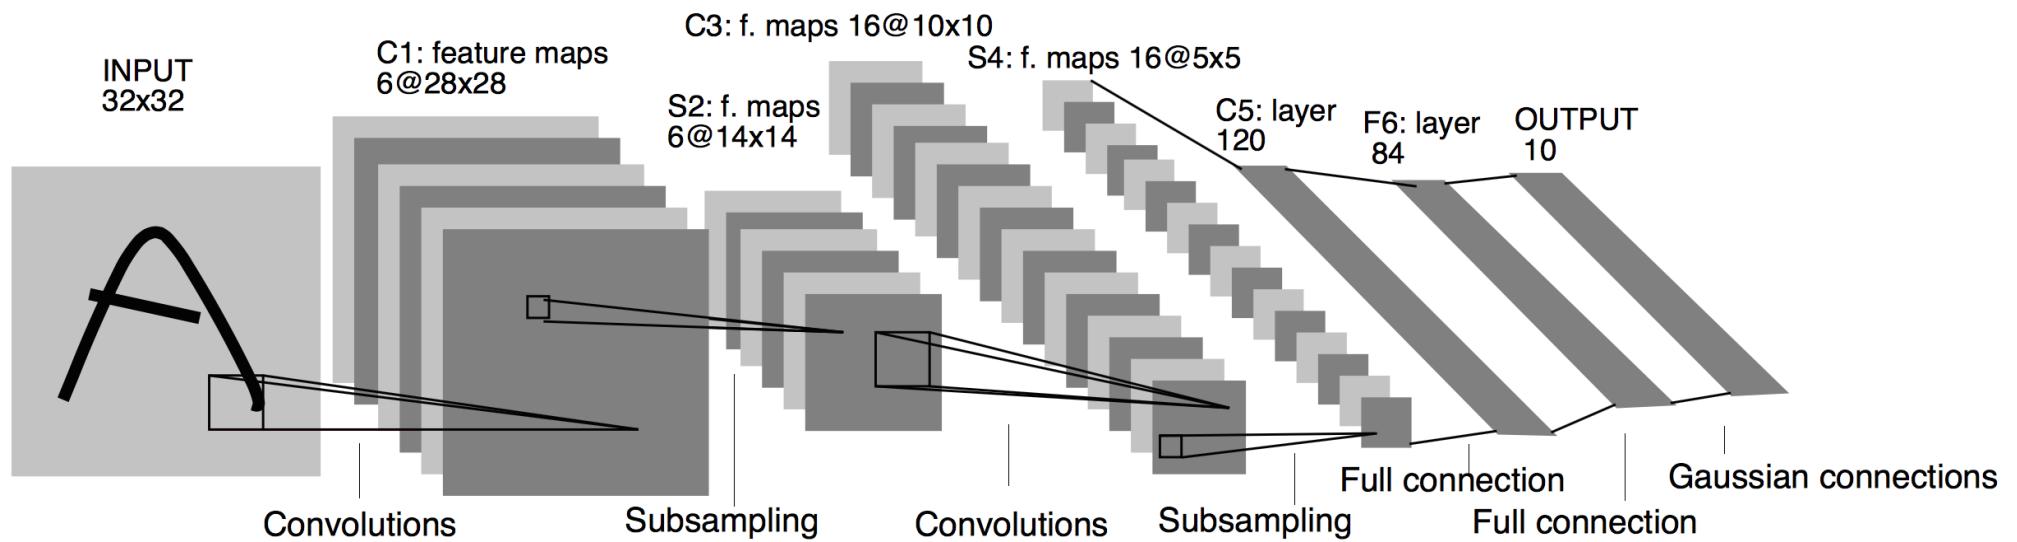
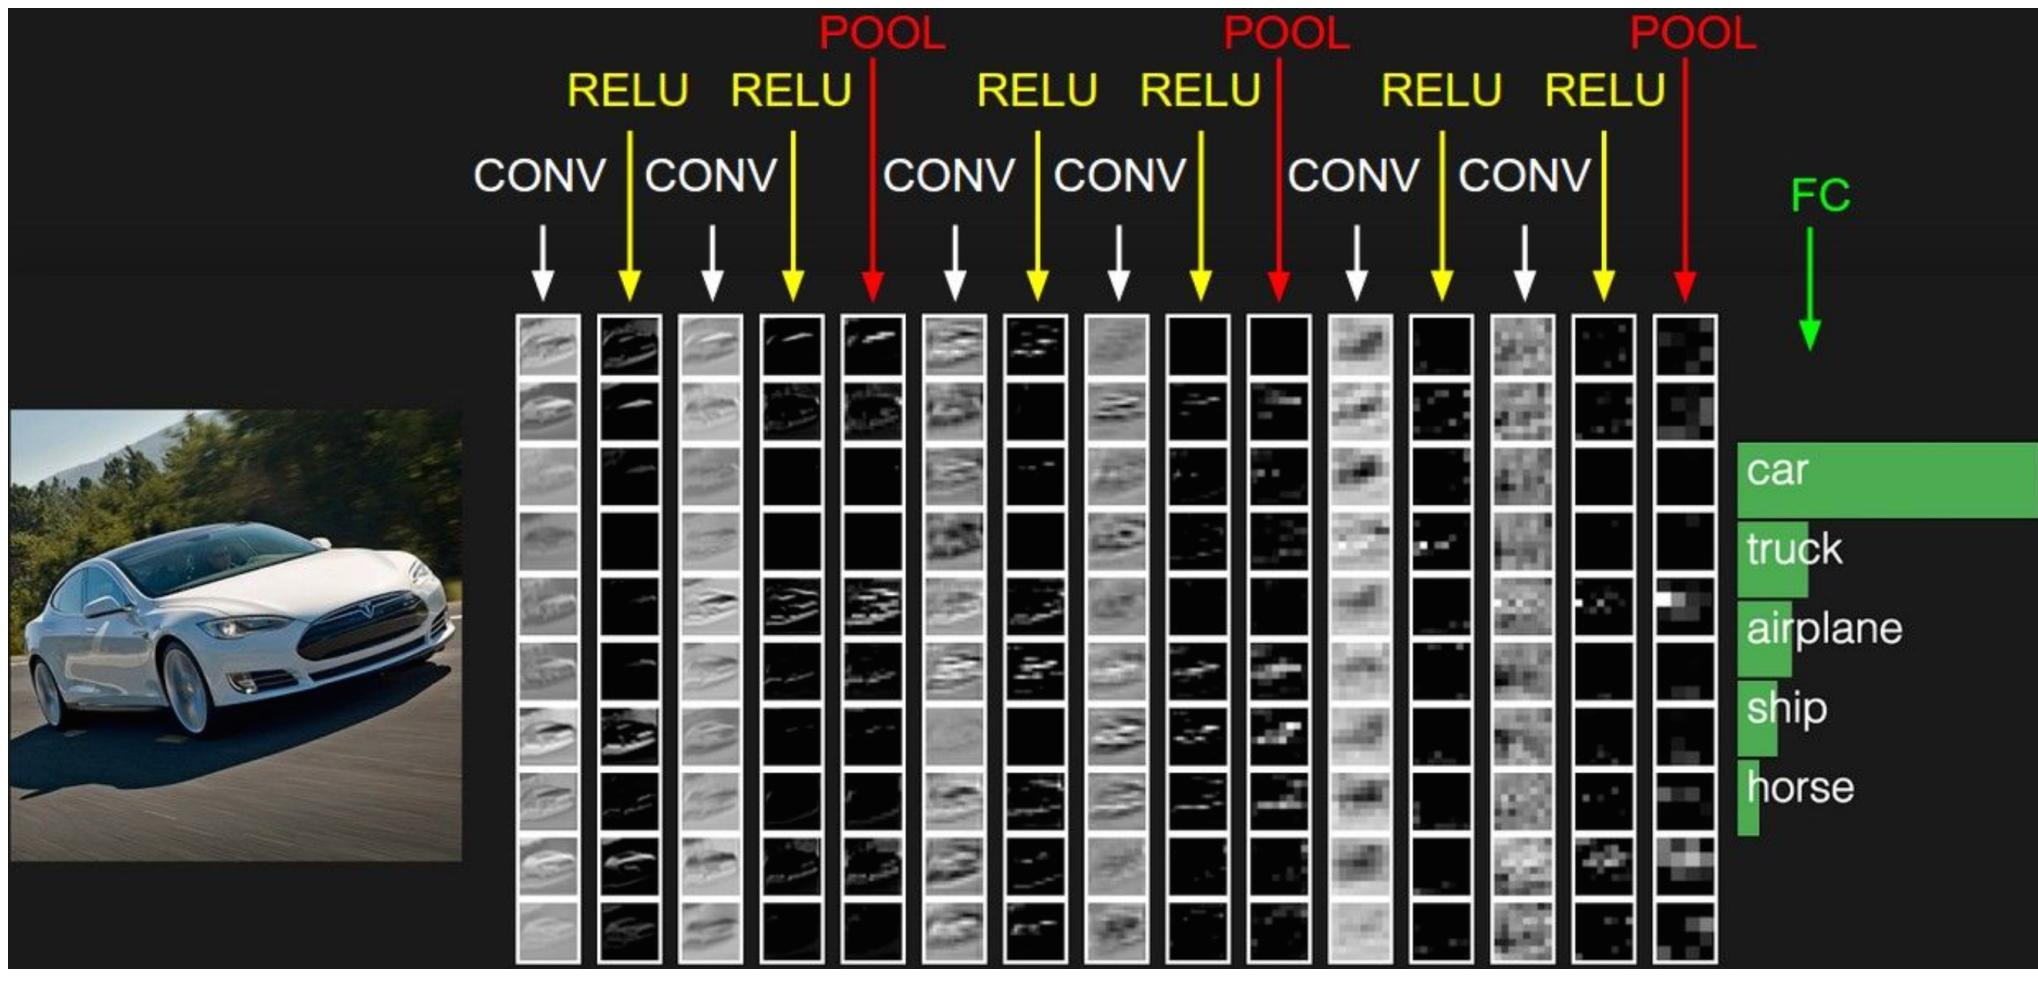


Fig. 2. Architecture of LeNet-5, a Convolutional Neural Network, here for digits recognition. Each plane is a feature map, i.e. a set of units whose weights are constrained to be identical.

Proc. Of the IEEE, November 1998, “Gradient-Based Learning Applied to Document Recognition”



Review: What is convolution?

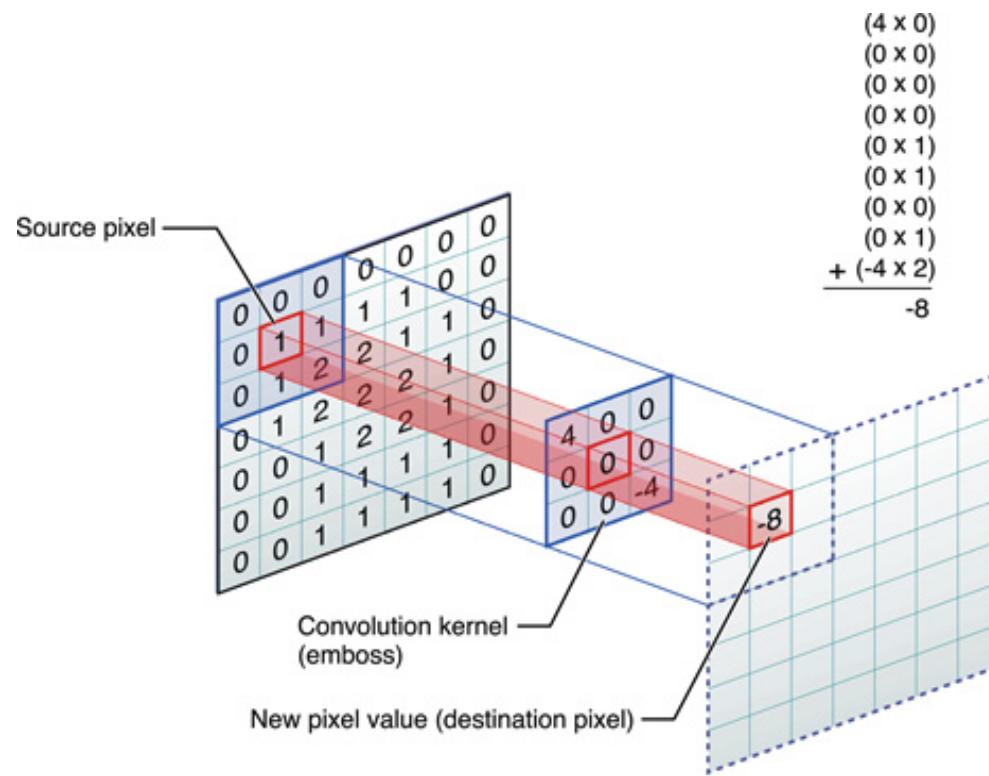
- Convolution is an important operation from signal processing
- A convolution is an integral that expresses the amount of overlap of one function as it is shifted over another function .

$$f * g = \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau = \int_{-\infty}^{\infty} g(\tau) f(t - \tau) d\tau$$

- 2 Dimensional Discrete Function (Image)

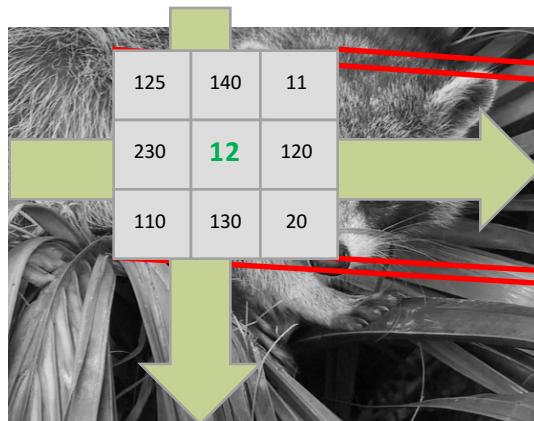
$$f[x, y] * g[x, y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1, n_2] \cdot g[x - n_1, y - n_2]$$

2-Dimensional Convolution

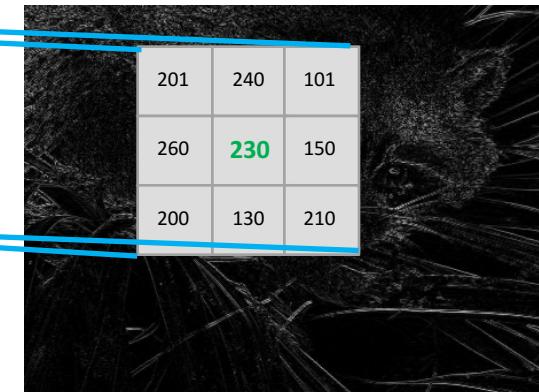


Example: 2-Dimensional Convolution

A convolution is an integral (**discrete signals :Matrix Dot Product**) that expresses the amount of overlap of one function as it is shifted over another function



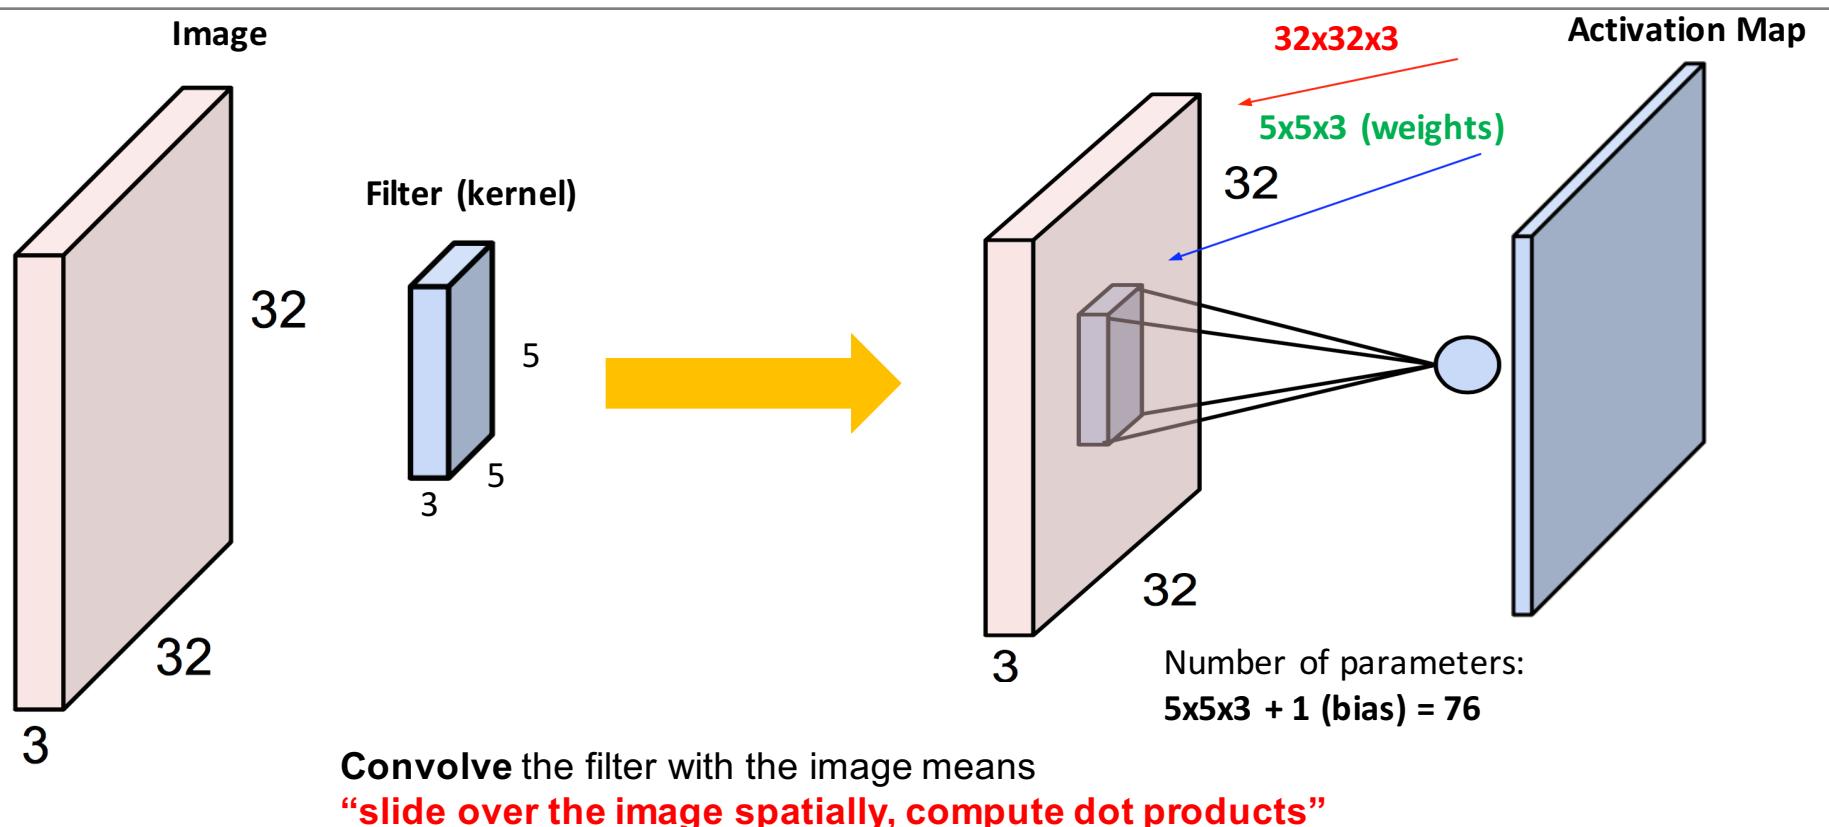
| | | |
|----|----|----|
| -1 | -2 | -1 |
| 0 | 0 | 0 |
| 1 | 2 | 1 |



| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |



Convolution Layer



Input Volume (+pad 1) (7x7x3)

| x[:, :, 0] |
|---------------|
| 0 0 0 0 0 0 0 |
| 0 2 0 1 1 0 0 |
| 0 2 1 2 2 1 0 |
| 0 2 0 0 1 2 0 |
| 0 1 1 2 2 1 0 |
| 0 0 1 0 2 2 0 |
| 0 0 0 0 0 0 0 |
| x[:, :, 1] |
| 0 0 0 0 0 0 0 |
| 0 1 2 1 1 2 0 |
| 0 1 2 1 2 0 0 |
| 0 2 0 1 2 2 0 |
| 0 2 2 2 1 0 0 |
| 0 0 1 0 2 2 0 |
| 0 0 0 0 0 0 0 |
| x[:, :, 2] |
| 0 0 0 0 0 0 0 |
| 0 0 0 2 0 0 0 |
| 0 1 1 1 0 2 0 |
| 0 2 1 1 2 1 0 |
| 0 0 2 1 1 0 0 |
| 0 0 0 2 1 2 0 |
| 0 0 0 0 0 0 0 |

Filter W0 (3x3x3)

| w0[:, :, 0] |
|-------------|
| 1 1 1 |
| 1 1 1 |
| 0 -1 0 |
| 0 1 1 |
| -1 1 -1 |
| 1 0 1 |
| -1 1 0 |
| -1 1 1 |
| 1 1 -1 |
| 1 1 1 |
| 1 1 0 |

Filter W1 (3x3x3)

| w1[:, :, 0] |
|-------------|
| 1 1 0 |
| 0 0 -1 |
| 0 0 1 |
| -1 0 -1 |
| -1 1 -1 |
| -1 0 1 |
| 1 1 -1 |
| -1 1 1 |
| 1 1 0 |

Output Volume

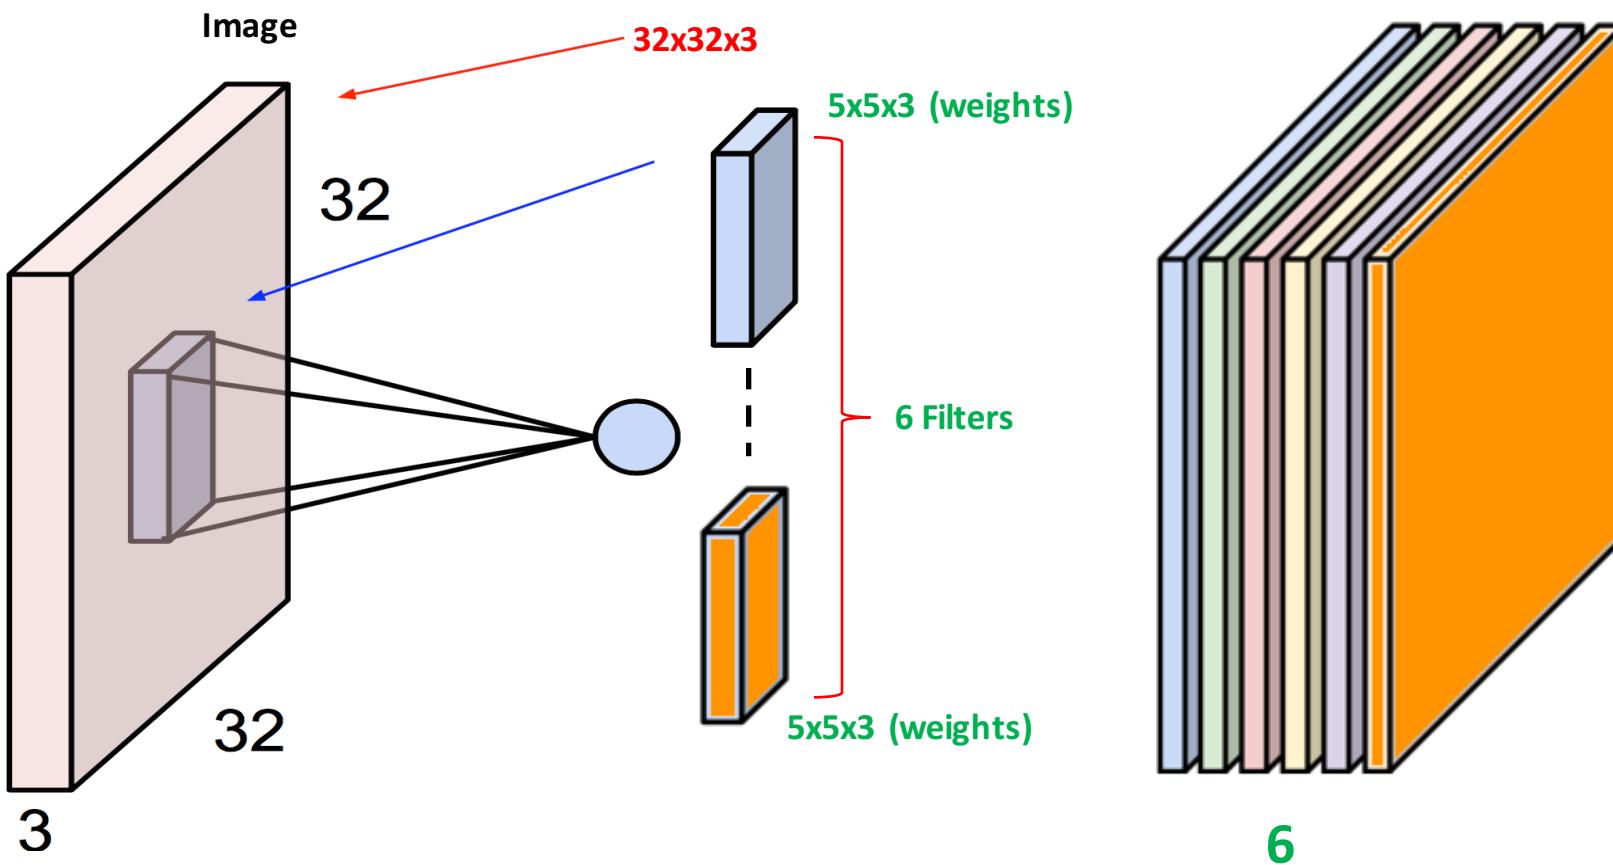
| o[:, :, 0] |
|------------|
| 9 |
| 2 |

$$(2 \times 1) + (1 \times 1) + 0 + (1 \times 1) + (2 \times 1) + 0 + (2 \times 1) + (1 \times 1) + 0 = 9$$

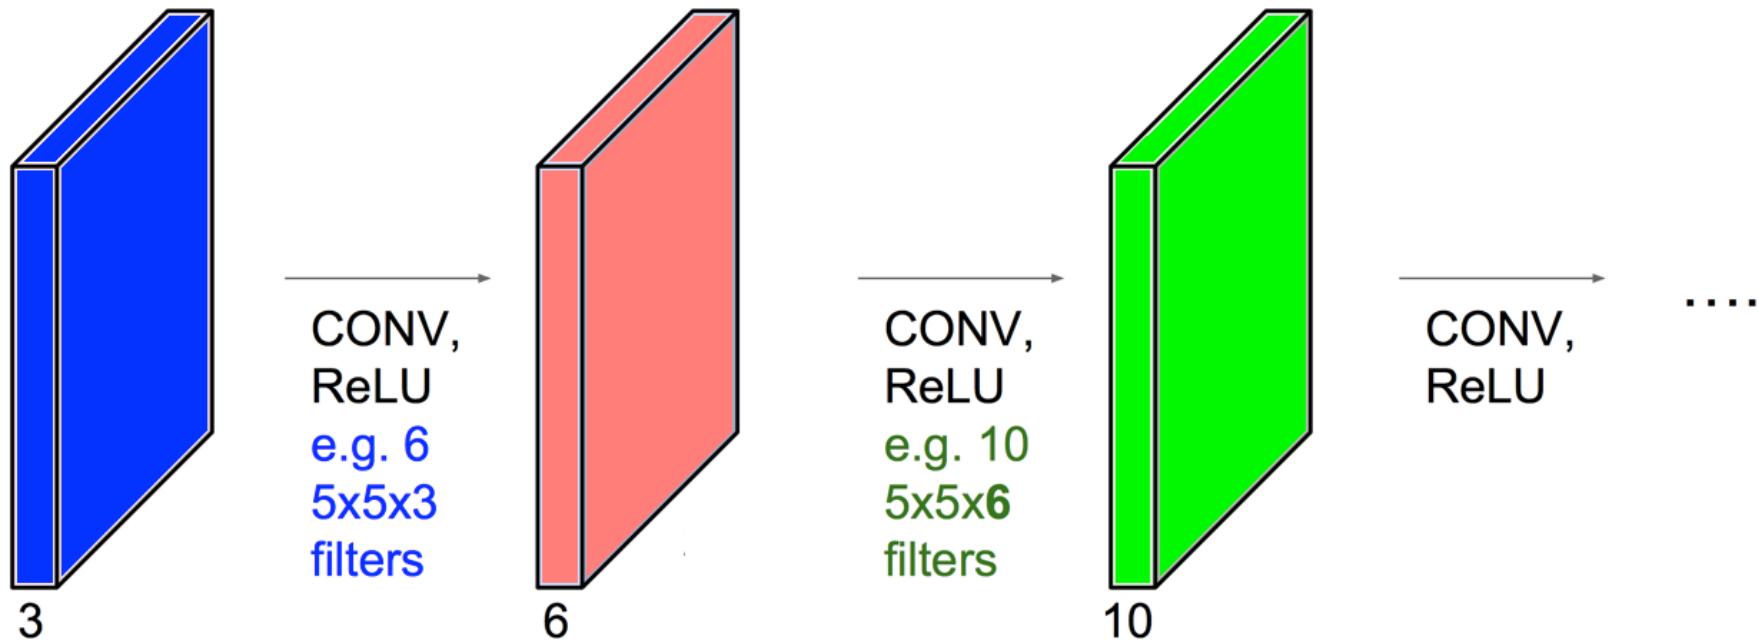
$$0 + 0 + 0 + (2 \times 1) + 0 + (1 \times -1) + 0 + 0 = 1$$

$$0 + 0 + 0 + (2 \times -1) + (1 \times 1) + 0 + (1 \times -1) + 0 + 0 = -2$$

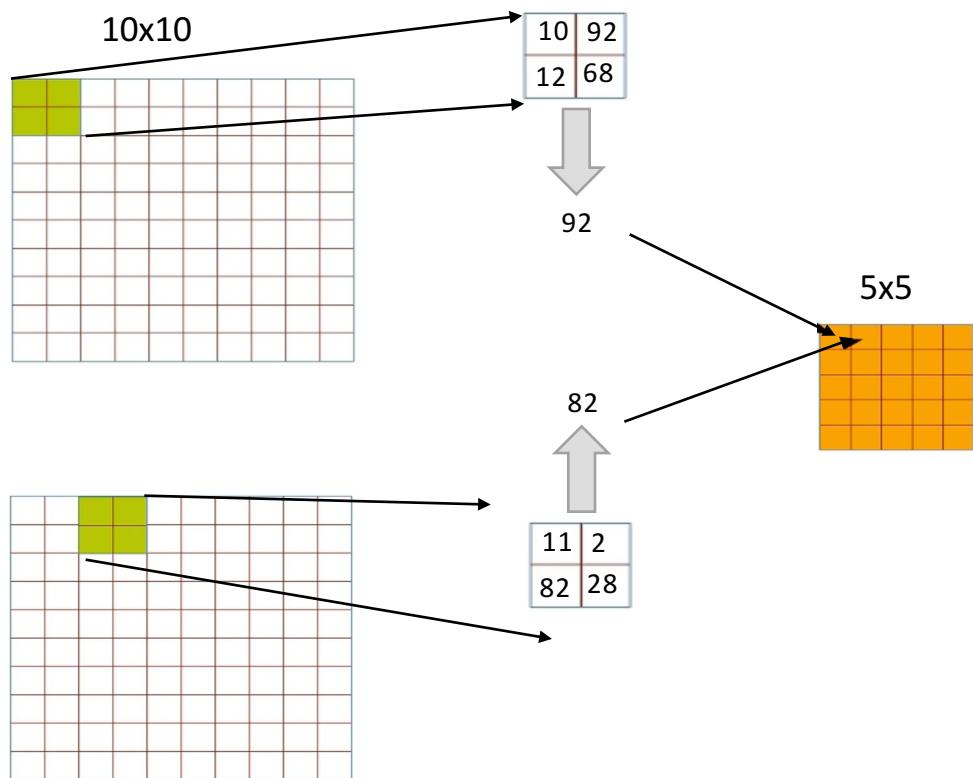
$$1 + 1 + (-2) + 9 = 9$$



Convolution Network



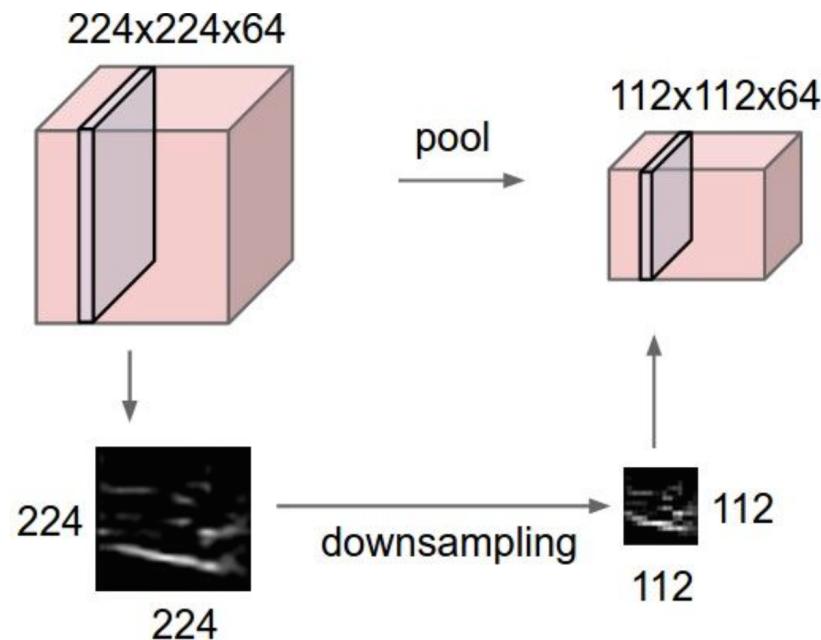
Downsampling= Max Pooling

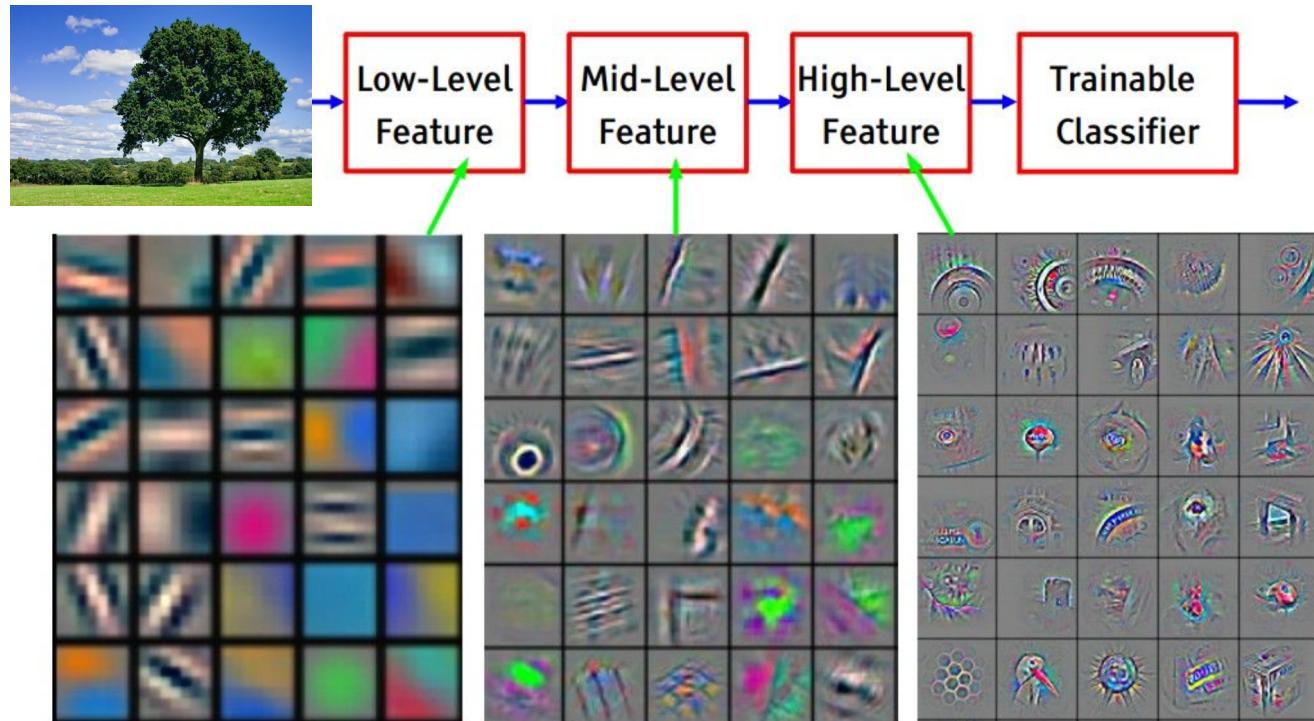


**“Max Pooling” Layers to extract the
“best” local feature**

Pooling Layer

- Makes the representations smaller and more manageable
- Operates over each activation map independently





Feature visualization of convolution net trained on ImageNet from [Zeiler & Fergus 2013]

